

# How to bind Dockers to NUMA nodes

## Overview

In this document I will introduce how to run the database dockers, e.g. redis, memcached, mongodb, on the specified NUMA node. Procedures should be similar if you want to run other kinds of dockers.

## Step1. Change the docker - install `numactl`

Specify the dockers you want to bind to a specific NUMA node. For example, I want to bind the redis server, memcached, and mongodb to the remote node. We should first launch the docker in interactive mode to see if `numactl` tool is supported:

```
sudo docker run -it <image_id> /bin/bash
```

You can use the following command to see the image ID:

```
sudo docker ps
```

Check if `numactl` is installed. If not, install it.

```
numactl -H
# if numactl is not found
apt-get update
apt-get install numactl
```

After you made the change, use the following commands to exit the docker and commit the changes:

```
exit
sudo docker ps -a # to find the container ID
sudo docker commit [CONTAINER_ID] [new_image_name] # e.g. redis-numa
```

**Note:** There is a special case for `mongo` docker. You may not commit now because further changes will be made.

## Step2. Change the docker - find the entrypoint

After you make sure `numactl` is installed, we should find the entrypoint of this docker. As a practice, the first command dockers will run is usually like `docker-entrypoint.sh redis-server`. You should check [dockerhub](#) to check the docker command used for those dockers that are directly pulled from dockerhub. As an example, the entry command for `redis` docker is [here](#). For a better understanding of the difference between `ENTRYPOINT` and `CMD` commands, you can check this [answer](#).

The entry command you find will later be added to the `docker-compose.yml` file.

## Step3. Modify `docker-compose.yml`

This file resides in the root path of a specific benchmark, e.g. `socialNetwork/`. It is good practice to "save as" the file and work on its copy. Modify a service in the file from:

```
social-graph-redis:
  image: redis
  hostname: social-graph-redis
  restart: always
```

to:

```
social-graph-redis:
  image: redis-numa // the image you created
  hostname: social-graph-redis
  privileged: true // numactl requires root privilege
  entrypoint: numactl --cpunodebind=0 --membind=1 redis-server // the entry command for the database docker with the numactl as prefix
  restart: always
```

You might have noticed that we did not use the `docker-entrypoint.sh` to start the database application. This depends on the docker. For `redis` as an example, the shell script does not really do anything, thus it is fairly safe to ignore it. However, for `mongo` the shell script does a

lot of initialization jobs so it should not be ignored. In that case, you should modify the `docker-entrypoint.sh` file inside the docker to **implant** the `numactl --cpunodebind=0 --membind=1` command.

Note: I will skip the steps to change `docker-entrypoint.sh` here, but as a hint the `numactl` is part of the mongo docker already. The final entrypoint command is something like `docker-entrypoint.sh mongod`.

## Step4. Run the dockers

Since you might want to switch between local and remote nodes, a shell script to rewrite `docker-compose.yml` is helpful. That is, you can create two `yml` files like `docker-compose-local.yml` and `docker-compose-remote.yml`, then use `cp` command to substitute the real `docker-compose.yml`.

Use the following command to set up the dockers.

```
sudo docker-compose up -d
```

Check the status of the dockers to make sure they are running normally (no restarting) and are actually started by a `numactl` command.

```
sudo docker ps
```